

Residual Dominance Drives Last-Item Reliance in Causal Self-Attention for Sequential Recommendation (Supplementary Document)

Anonymous Author(s)

Important Notice

If the text on this page does not display correctly, reloading the page may resolve the issue. Please make sure you are viewing the most recent version.

1 Further Details on Norm-Based Analysis

In this appendix, we provide the full mathematical derivation of the norm-based analysis used in Section 4. We follow the framework of [2, 3], who extended the norm-based analysis of from multi-head attention alone to the entire attention block (multi-head attention, residual connection, and layer normalization). We describe how each component is decomposed and how the context-mixing ratio is computed at each stage.

1.1 Attention Block Formulation

We consider the attention block used in SASRec [1]. Given input representations $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n] \in \mathbb{R}^{n \times d}$, the output at position i is:

$$\tilde{\mathbf{x}}_i = \text{LN}(\text{Attn}(\mathbf{x}_i, \mathbf{X}) + \mathbf{x}_i), \quad (1)$$

where $\mathbf{x}_i \in \mathbb{R}^d$ denotes the i -th input representation (row vector), and $\tilde{\mathbf{x}}_i \in \mathbb{R}^d$ is the corresponding output. This block consists of three components: multi-head attention (Attn), residual connection, and layer normalization (LN).

Our goal is to decompose $\tilde{\mathbf{x}}_i$ into the sum of vector-valued functions of each input $\mathbf{x}_j$, and quantify their relative magnitudes.

1.2 Decomposition of Multi-Head Attention

The H -head attention output is:

$$\text{Attn}(\mathbf{x}_i, \mathbf{X}) = \sum_{h=1}^H \text{Attn}^{(h)}(\mathbf{x}_i, \mathbf{X}), \quad (2)$$

where each head computes:

$$\text{Attn}^{(h)}(\mathbf{x}_i, \mathbf{X}) = \sum_{j=1}^n \alpha_{i,j}^{(h)} \left(\mathbf{x}_j \mathbf{W}_V^{(h)} + \mathbf{b}_V^{(h)} \right) \mathbf{W}_O^{(h)}, \quad (3)$$

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Conference'17, Washington, DC, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-1-4503-XXXX-X/2026/06
<https://doi.org/XXXXXXX.XXXXXXX>

with attention weights:

$$\alpha_{i,j}^{(h)} = \text{softmax}_{\mathbf{x}_j \in \mathbf{X}} \left(\frac{(\mathbf{x}_i \mathbf{W}_Q^{(h)} + \mathbf{b}_Q^{(h)})(\mathbf{x}_j \mathbf{W}_K^{(h)} + \mathbf{b}_K^{(h)})^\top}{\sqrt{d_h}} \right). \quad (4)$$

Here, $\mathbf{W}_Q^{(h)}, \mathbf{W}_K^{(h)}, \mathbf{W}_V^{(h)} \in \mathbb{R}^{d \times d_h}$ and $\mathbf{W}_O^{(h)} \in \mathbb{R}^{d_h \times d}$ are weight matrices, and $\mathbf{b}_Q^{(h)}, \mathbf{b}_K^{(h)}, \mathbf{b}_V^{(h)} \in \mathbb{R}^{d_h}$ are bias vectors, with $d_h = d/H$. Since SASRec uses causal (left-to-right) self-attention, $\alpha_{i,j}^{(h)} = 0$ for $j > i$.

We define the value transformation for each head as:

$$f^{(h)}(\mathbf{x}) := (\mathbf{x} \mathbf{W}_V^{(h)} + \mathbf{b}_V^{(h)}) \mathbf{W}_O^{(h)}. \quad (5)$$

Substituting into (2) and (3), the attention output decomposes into a sum over input positions:

$$\text{Attn}(\mathbf{x}_i, \mathbf{X}) = \sum_{j=1}^n \sum_{h=1}^H \alpha_{i,j}^{(h)} f^{(h)}(\mathbf{x}_j). \quad (6)$$

This decomposition, introduced by Kobayashi et al. [2], shows that the attention output is naturally a sum of transformed input vectors weighted by attention.

1.3 Incorporating the Residual Connection

Adding the residual connection to (6):

$$\mathbf{y}_i := \text{Attn}(\mathbf{x}_i, \mathbf{X}) + \mathbf{x}_i = \sum_{h=1}^H \sum_{j=1}^n \alpha_{i,j}^{(h)} f^{(h)}(\mathbf{x}_j) + \mathbf{x}_i. \quad (7)$$

Since the residual term $\mathbf{x}_i$ involves only position i , the additive structure is preserved. We separate the right-hand side into per-position contributions $\mathbf{y}_{i,j}$ as follows.

For $j \neq i$ (contributions from context):

$$\mathbf{y}_{i,j} := \sum_{h=1}^H \alpha_{i,j}^{(h)} f^{(h)}(\mathbf{x}_j). \quad (8)$$

For $j = i$ (contribution from self, via both attention and residual):

$$\mathbf{y}_{i,i} := \sum_{h=1}^H \alpha_{i,i}^{(h)} f^{(h)}(\mathbf{x}_i) + \mathbf{x}_i. \quad (9)$$

By construction, $\mathbf{y}_i = \sum_j \mathbf{y}_{i,j}$.

1.4 Decomposition of Layer Normalization

Layer normalization is defined as:

$$\text{LN}(\mathbf{y}) = \frac{\mathbf{y} - m(\mathbf{y})}{s(\mathbf{y})} \odot \boldsymbol{\gamma} + \boldsymbol{\beta}, \quad (10)$$

where $\odot$ denotes element-wise multiplication, and

$$m(\mathbf{y}) = \frac{1}{d} \sum_{k=1}^d y^{(k)}, \quad s(\mathbf{y}) = \sqrt{\frac{1}{d} \sum_{k=1}^d (y^{(k)} - m(\mathbf{y}))^2 + \epsilon}, \quad (11)$$

with $y^{(k)}$ denoting the k -th element of $\mathbf{y}$, and $\epsilon > 0$ a small constant for numerical stability. Here, $\boldsymbol{\gamma}, \boldsymbol{\beta} \in \mathbb{R}^d$ are learnable parameters for element-wise scaling and shifting, respectively.

We now show that LN admits a “distributive law” over the additive decomposition of its input. Let $\mathbf{y}_i = \sum_j \mathbf{y}_{i,j}$ be the input to LN. Substituting into (10):

$$\text{LN}(\mathbf{y}_i) = \frac{\mathbf{y}_i - m(\mathbf{y}_i)}{s(\mathbf{y}_i)} \odot \boldsymbol{\gamma} + \boldsymbol{\beta} = \frac{\sum_j \mathbf{y}_{i,j} - m(\sum_j \mathbf{y}_{i,j})}{s(\mathbf{y}_i)} \odot \boldsymbol{\gamma} + \boldsymbol{\beta}. \quad (12)$$

The mean operator $m(\cdot)$ is a linear function (it computes the arithmetic mean of the vector elements). Therefore:

$$m\left(\sum_j \mathbf{y}_{i,j}\right) = \sum_j m(\mathbf{y}_{i,j}). \quad (13)$$

Substituting (13) into (12):

$$\text{LN}(\mathbf{y}_i) = \frac{\sum_j \mathbf{y}_{i,j} - \sum_j m(\mathbf{y}_{i,j})}{s(\mathbf{y}_i)} \odot \boldsymbol{\gamma} + \boldsymbol{\beta} \quad (14)$$

$$= \frac{\sum_j (\mathbf{y}_{i,j} - m(\mathbf{y}_{i,j}))}{s(\mathbf{y}_i)} \odot \boldsymbol{\gamma} + \boldsymbol{\beta} \quad (15)$$

$$= \sum_j \frac{\mathbf{y}_{i,j} - m(\mathbf{y}_{i,j})}{s(\mathbf{y}_i)} \odot \boldsymbol{\gamma} + \boldsymbol{\beta}. \quad (16)$$

In the last step, we use the fact that the denominator $s(\mathbf{y}_i)$ is a scalar computed from the *entire* input $\mathbf{y}_i = \sum_j \mathbf{y}_{i,j}$, and is therefore a constant with respect to the summation over j . This allows it to be factored out of the sum. Note that, unlike $m(\cdot)$, the standard deviation $s(\cdot)$ is a nonlinear function and cannot be distributed over the sum. Instead, $s(\mathbf{y}_i)$ serves as a shared normalization factor across all components $\mathbf{y}_{i,j}$.

Defining:

$$g_{\mathbf{y}_i}(\mathbf{y}_{i,j}) := \frac{\mathbf{y}_{i,j} - m(\mathbf{y}_{i,j})}{s(\mathbf{y}_i)} \odot \boldsymbol{\gamma}, \quad (17)$$

we obtain the decomposition:

$$\tilde{\mathbf{x}}_i = \text{LN}(\mathbf{y}_i) = \sum_j g_{\mathbf{y}_i}(\mathbf{y}_{i,j}) + \boldsymbol{\beta}. \quad (18)$$

This decomposition is *exact*; no approximation is involved. The numerator of each term $g_{\mathbf{y}_i}(\mathbf{y}_{i,j})$ is specific to the component $\mathbf{y}_{i,j}$ (the mean $m(\mathbf{y}_{i,j})$ is computed from $\mathbf{y}_{i,j}$ alone), while the denominator $s(\mathbf{y}_i)$ is shared across all components.

1.5 Full Decomposition

Combining the decompositions of Attn (§6), residual connection (§8–9), and LN (§18), the output of the attention block is:

$$\tilde{\mathbf{x}}_i = \underbrace{\sum_{j \neq i} g_{\mathbf{y}_i}(\mathbf{y}_{i,j})}_{\text{context-mixing}} + \underbrace{g_{\mathbf{y}_i}(\mathbf{y}_{i,i})}_{\text{self-preserving}} + \boldsymbol{\beta}, \quad (19)$$

where $\mathbf{y}_{i,j}$ is given by (8) for $j \neq i$ and by (9) for $j = i$, and $g_{\mathbf{y}_i}$ is defined in (17).

The *context-mixing* terms capture information flowing from other positions into the output at position i , while the *self-preserving* term captures how much of the original input $\mathbf{x}_i$ is retained. In the self-preserving term, information from $\mathbf{x}_i$ flows through two pathways: self-attention (weighted by $\alpha_{i,i}^{(h)}$) and the residual connection.

1.6 Context-Mixing Ratio

Based on the decomposition in (19), we define the context-mixing and self-preserving vectors as:

$$\tilde{\mathbf{x}}_{i \leftarrow \text{context}} := \sum_{j \neq i} g_{\mathbf{y}_i}(\mathbf{y}_{i,j}), \quad (20)$$

$$\tilde{\mathbf{x}}_{i \leftarrow i} := g_{\mathbf{y}_i}(\mathbf{y}_{i,i}). \quad (21)$$

The **context-mixing ratio** is defined as:

$$r_i = \frac{\|\tilde{\mathbf{x}}_{i \leftarrow \text{context}}\|}{\|\tilde{\mathbf{x}}_{i \leftarrow \text{context}}\| + \|\tilde{\mathbf{x}}_{i \leftarrow i}\|}. \quad (22)$$

A higher r_i indicates that context-mixing is dominant in computing $\tilde{\mathbf{x}}_i$; a lower r_i indicates that self-preserving is dominant. The bias term $\boldsymbol{\beta}$ affects neither the context-mixing nor the self-preserving effect, and is therefore excluded from the analysis.

Note that r_i is defined as the ratio of the *norm of the vector sum*, not the sum of individual norms. That is, we first aggregate the contribution vectors, then measure the magnitude of the aggregated vector.

1.7 Stage-Wise Mixing Ratios

To isolate the role of each component in the attention block, we compute mixing ratios at three stages, following the analysis framework of Kobayashi et al. [2].

1.7.1 Attention Only (Attn-N). Considering only the multi-head attention output (6), without residual connection or layer normalization:

$$\mathbf{c}_i^{\text{Attn}} := \sum_{j \neq i} \sum_{h=1}^H \alpha_{i,j}^{(h)} f^{(h)}(\mathbf{x}_j), \quad (23)$$

$$\mathbf{p}_i^{\text{Attn}} := \sum_{h=1}^H \alpha_{i,i}^{(h)} f^{(h)}(\mathbf{x}_i), \quad (24)$$

$$r_i^{\text{Attn}} := \frac{\|\mathbf{c}_i^{\text{Attn}}\|}{\|\mathbf{c}_i^{\text{Attn}}\| + \|\mathbf{p}_i^{\text{Attn}}\|}. \quad (25)$$

1.7.2 Attention + Residual (AttnRes-N). Incorporating the residual connection (8)–(9), but without layer normalization:

$$\mathbf{c}_i^{\text{AttnRes}} := \sum_{j \neq i} \sum_{h=1}^H \alpha_{i,j}^{(h)} f^{(h)}(\mathbf{x}_j), \quad (26)$$

$$\mathbf{p}_i^{\text{AttnRes}} := \sum_{h=1}^H \alpha_{i,i}^{(h)} f^{(h)}(\mathbf{x}_i) + \mathbf{x}_i, \quad (27)$$

$$r_i^{\text{AttnRes}} := \frac{\|\mathbf{c}_i^{\text{AttnRes}}\|}{\|\mathbf{c}_i^{\text{AttnRes}}\| + \|\mathbf{p}_i^{\text{AttnRes}}\|}. \quad (28)$$

Note that $\mathbf{c}_i^{\text{AttnRes}} = \mathbf{c}_i^{\text{Attn}}$, since the residual connection adds $\mathbf{x}_i$ only to the self-preserving term. The residual connection thus increases $\|\mathbf{p}_i^{\text{AttnRes}}\|$ relative to $\|\mathbf{p}_i^{\text{Attn}}\|$, which in general decreases the mixing ratio.

1.7.3 *Full Attention Block (AttnResLN-N)*. Incorporating all three components, using the decomposition derived in (19):

$$\mathbf{c}_i^{\text{AttnResLN}} := \sum_{j \neq i} g_{y_i}(\mathbf{y}_{i,j}), \quad (29)$$

$$\mathbf{p}_i^{\text{AttnResLN}} := g_{y_i}(\mathbf{y}_{i,i}), \quad (30)$$

$$\mathbf{r}_i^{\text{AttnResLN}} := \frac{\|\mathbf{c}_i^{\text{AttnResLN}}\|}{\|\mathbf{c}_i^{\text{AttnResLN}}\| + \|\mathbf{p}_i^{\text{AttnResLN}}\|}. \quad (31)$$

Kobayashi et al. [3] showed that layer normalization further reduces the context-mixing ratio by scaling (via γ) in a way that shrinks the attention-derived vectors relative to the residual.

References

- [1] Wang-Cheng Kang and Julian McAuley. 2018. Self-attentive sequential recommendation. In *Proceedings of the IEEE International Conference on Data Mining*. 197–206.
- [2] Goro Kobayashi, Tatsuki Kuribayashi, Sho Yokoi, and Kentaro Inui. 2020. Attention is not only a weight: Analyzing transformers with vector norms. *arXiv preprint arXiv:2004.10102* (2020).
- [3] Goro Kobayashi, Tatsuki Kuribayashi, Sho Yokoi, and Kentaro Inui. 2021. Incorporating residual and normalization layers into analysis of masked language models. *arXiv preprint arXiv:2109.07152* (2021).